

Monica Kodwani

PHD STUDENT · COMPUTER SCIENCE (USABLE SECURITY & PRIVACY; TRUSTWORTHY AI)

Washington, D.C., USA

☎ (+1) 202-428-8379 | ✉ monicakodwani@gwu.edu | 🌐 www.monicakodwani.com | 📱 monicakodwani

Research Areas: Trustworthy AI | Privacy, Safety, Security | Human-Centered Design | Fairness and Accountability

Education

George Washington University

PH.D. IN COMPUTER SCIENCE (USABLE SECURITY & PRIVACY), EXPECTED GRADUATION: MAY 2028

- Advisor: Dr. Adam Aviv

Washington, D.C., USA

May 2023 – Present

George Washington University

MASTER OF SCIENCE IN CYBERSECURITY IN COMPUTER SCIENCE

Washington, D.C., USA

Jan 2022 – May 2023

Manipal University Jaipur

BACHELOR OF TECHNOLOGY IN COMPUTER COMMUNICATION ENGINEERING

Jaipur, India

Aug 2016 – May 2020

Publications

How Probing for Problems and Bias Affects Perceptions of AI Chatbot Trustworthiness

Conference: ACM Conference on Fairness, Accountability, and Transparency (FAccT 2026) - To Appear

Authors: Jan Tolsdorf, Alan F. Luo, **Monica Kodwani**, Junho Eum, Mahmood Sharif, Michelle L. Mazurek, and Adam J. Aviv

Abstract: Conducted a large-scale interactive study (N=254, 551 conversations) to investigate how lay users evaluate the trustworthiness of generative AI systems. Found that users often fail to surface expected issues and may increase trust even after limited probing, revealing systematic gaps in user-driven auditing. This work informs the design of human-centered evaluation frameworks and interventions for trustworthy AI.

Safety Perceptions of Generative AI Conversational Agents: Uncovering Perceptual Differences in Trust, Risk, and Fairness

Conference: Proceedings of the 21st Symposium on Usable Privacy and Security (SOUPS) 2025

Authors: Jan Tolsdorf, Alan F. Luo, **Monica Kodwani**, Junho Eum, Mahmood Sharif, Michelle L. Mazurek, and Adam J. Aviv

Abstract: Surveyed 123 U.S. chatbot users with a new trust-risk-fairness instrument. Factor and latent-class analyses revealed three user segments—Hesitant Skeptics, Cautious Trusters, Confident Adopters—showing usage frequency, not demographics, is the strongest predictor of higher trust and lower perceived risk.

On a Scale of 1 to 5, How Reliable Are AI User Studies? A Call for Developing Validated, Meaningful Scales and Metrics about User Perceptions of AI Systems.

Conference: IEEE SPW Workshop on Technology and Consumer Protection (ConPro) 2025

Authors: Jan Tolsdorf, Alan F. Luo, **Monica Kodwani**, Junho Eum, Mahmood Sharif, Michelle L. Mazurek, and Adam J. Aviv

Abstract: Literature review of 130 papers shows most AI-perception studies recycle unvalidated Likert items. Calls for a community repository of rigorously tested scales and outlines a roadmap for scale co-development, piloting, and open benchmarking.

Honesty is the Best Policy: On the Accuracy of Apple Privacy Labels Compared to Apps' Privacy Policies

Conference: Proceedings on Privacy Enhancing Technologies (PETS) 2024

Authors: Mir Masood Ali, David G. Balash, **Monica Kodwani**, Chris Kanich, Adam J. Aviv

Abstract: Conducted a large-scale analysis comparing Apple's privacy labels with actual app privacy policies using fine-tuned BERT-based language models. Identified significant discrepancies, providing insights to improve the accuracy of privacy labels and assist developers.

Longitudinal Analysis of Privacy Labels in the Apple App Store

arXiv:2206.02658

Authors: David G. Balash, Mir Masood Ali, **Monica Kodwani**, Xiaoyuan Wu, Chris Kanich, Adam J. Aviv

Abstract: Collected weekly snapshots of 1.6M apps (Jul 2021–Oct 2022) and found 70% label adoption—driven largely by new apps—while 42% of labeled apps claim “Data Not Collected.” Labels changed rarely and usually expanded data collection, raising questions about reliability.

Research Experience

George Washington University

Washington, D.C., USA

RESEARCH ASSISTANT

May 2023 – Present

- **Human-Centered Evaluation and Auditing of Generative AI**
 - Led a 250-participant study using a custom interactive audit platform to examine how users identify bias, fairness concerns, and other problematic behaviors in chatbot interactions.
 - Designed and deployed a chat-based survey and annotation system supporting structured auditing workflows for generative AI.
- **Retrospective Privacy and AI Interaction History**
 - Designed a dashboard-based study investigating how reviewing long-term AI interaction history affects users’ awareness of sensitive information disclosure, privacy concerns, and trust.
- **Social Bias in Generative AI**
 - Initiated research on caste bias in image generation models, examining how social hierarchies and representation harms manifest in generated visual outputs.
- **AI in Workplace Contexts**
 - Collaborating with NIST on an interview study examining AI use, trust, and governance in workplace settings.
- **Collaboration & Leadership:** Mentored 2 PhD + 4 UG researchers; Collaborated with Principal Investigators from University of Maryland, University of Illinois Urbana-Champaign, CISPA Helmholtz Center for Information Security, & National Institute of Standards and Technology, US Department of Commerce

Technical Skills

Programming Languages	• Python (7 yrs), Java, JavaScript, C/C++, C#, Bash.
ML/NLP	• PyTorch, TensorFlow, HuggingFace, scikit-learn; Prompt Engineering & Fine Tuning LLMs (LoRA, QLoRA)
Data Engineering	• Elasticsearch clusters (sharded, ILM policies), PostgreSQL, AWS (S3, EC2, Lambda, SageMaker), Docker/K8s
Human-Centered Methods	• Online controlled experiments, Interviews (Semi-Structured), Survey Design (Qualtrics, Custom React), Psychometrics (EFA/CFA, Measurement Invariance), IRB/Ethics Compliance
Web & Visualization	• React 18, Next.js, D3.js, Plotly; Micro-Frontend Architecture.
Dev & Ops	• Git/GitHub Actions CI/CD, pytest, Terraform, Agile/Scrum Facilitation.

Selected Projects

PRIVACY-LABEL OBSERVATORY	2023-24
• Built a Next.js+Elasticsearch site that visualizes 3M-plus historical Apple/Play labels • AWS Batch crawlers refresh data weekly and drive interactive per-app timelines used by researchers studying mobile-privacy trends	
CUSTOM CHAT-SURVEY PLATFORM	2024
• Multi-page React/Node tool combining LLM chat, pre/post surveys, and highlight-and-annotate auditing workflow. • Deployed in a 250-participant survey to capture bias perceptions in-situ; supports Likert ratings and open-ended rationales.	
POLICY-LABEL NLP BENCHMARK	2024
• Fine-tuned transformer models to compare stated privacy labels with policy text for “Honesty is the Best Policy” project; benchmark evaluates mis-alignment detection accuracy, laying groundwork for automated label auditing.	

Work Experience

George Washington University

Washington, D.C., USA

GRADER – COMPUTER SECURITY COURSE

Fall 2022, Spring 2023, Fall 2024

- Assessed programming labs & exams for 60 students/semester; authored new buffer-overflow assignment, & hacking exercise adopted by course coordinator.

George Washington University

Washington, D.C., USA

IT SUPPORT STAFF

Spring 2022

- Resolved 98% of 200+ tickets within SLA; automated imaging script that cut lab PC deployment time -50%.

National Critical Infrastructure & Information Protection Center (NCIIPC), India

India

SECURITY ANALYSIS INTERN

2020

- Developed ZeroNet crawler monitoring 2K+ sites, alerting SOC within <5 min of malicious signature detection (20% reduction in incident rate).
- Implemented Onion-link extractor, lowering threat-triage time by 40%.

References

Available upon request.